

AUP : Assignment - 7 [Signals]

Aditi Rajendra Medhane 111803177

18th October 2021

Q1

Create a child process. Let the parent sleeps of 5 seconds and exits. Can the child send SIGINT to its parent if exists and kill it? Verify with a sample program.

Code

```
1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <sys/types.h>
4  #include <unistd.h>
5  #include <signal.h>
6  #include <errno.h>
7
8  int main(void) {
9      int child_pid;
10
11      if ((child_pid = fork()) == -1) {
12          //fork failed
13          perror("fork");
14          return errno;
15      }
16      else if (child_pid) {
17          //parent
18          sleep(5);
19          printf("Parent is Alive\n");
20      }
21      else {
22          //child
23          if (kill(getppid(), SIGINT) == -1) {
24              perror("SIGINT to parent");
25              return errno;
26          }
27          printf("Child sent kill signal\n");
28          sleep(5);
29          printf("Child slept for time equal to parent\n");
30      }
31
32      return 0;
33 }
```

Explanation

Parent is alive not printed, parent killed by SIGINT sent by child

Output

```
hp@aditi:~/Desktop/BTech/AUP/LAB/7$ ./1
Child sent kill signal

hp@aditi:~/Desktop/BTech/AUP/LAB/7$ Child slept for time equal to parent
```

Figure 1: Output 1

Q2

Create a signal disposition to catch SIGCHLD and in the handler function display some message. Create a child process and let the child sleeps for some time and exits. The parent calls a wait() for the child. Display the return value of wait() to check success or failure. If failure, display the error number. Run the program:

- Normal way executing in the foreground
- Run as a back ground process and send SIGCHLD to it from the shell

Code

```
1  #include <sys/types.h>
2  #include <sys/wait.h>
3  #include <unistd.h>
4  #include <signal.h>
5  #include <errno.h>
6  #include <stdio.h>
7
8  void print_sigchld_msg(int sig_number) {
9      printf("SIGCHLD hits parent (msg)\n");
10 }
11
12 int main(void) {
13     pid_t child_id;
14
15     int dead_child;
16     int status;
17
18     /* set handler for SIGCHLD */
19     signal(SIGCHLD, print_sigchld_msg);
20
21     if ((child_id = fork()) == -1) {
22         perror("fork");
23         return errno;
24     }
25     else if (child_id) {
26         /* parent */
27
28         if ((dead_child = wait(&status)) == -1) {
29             perror("wait failed");
```

```

30         return errno;
31     }
32     else {
33         printf("Wait returned, child %d exited\n", dead_child);
34     }
35 }
36 else {
37     /* Child */
38     printf("Child sleeping\n");
39     sleep(10);
40 }
41
42 return 0;
43 }

```

Explanation

- Image 2_1 : Normal way of executing the program, SIGCHLD hits once
- Image 2_2 : Run as a back ground process and send SIGCHLD to it from the shell, SIGCHLD hits twice

Output

```

hp@aditi:~/Desktop/BTech/AUP/LAB/7$ ./2
Child sleeping
SIGCHLD hits parent (msg)
Wait returned, child 58150 exited
hp@aditi:~/Desktop/BTech/AUP/LAB/7$

```

```

hp@aditi:~/Desktop/BTech/AUP/LAB/7$ ./2
Child sleeping
SIGCHLD hits parent (msg)
Wait returned, child 58334 exited
hp@aditi:~/Desktop/BTech/AUP/LAB/7$ ./2 &
[1] 58336
hp@aditi:~/Desktop/BTech/AUP/LAB/7$ Child sleeping
ps
  PID TTY          TIME CMD
 57904 pts/1    00:00:00 bash
 58336 pts/1    00:00:00 2
 58337 pts/1    00:00:00 2
 58338 pts/1    00:00:00 ps
hp@aditi:~/Desktop/BTech/AUP/LAB/7$ kill -17 58336
hp@aditi:~/Desktop/BTech/AUP/LAB/7$ SIGCHLD hits parent (msg)
SIGCHLD hits parent (msg)
Wait returned, child 58337 exited

```

Q3

You have to create a process tree as shown below. Then you let the parent process create a process group of (3, 4, 5) so that it sends a signal to this group. Print appropriate messages.

Code

```
1  #include <sys/types.h>
2  #include <sys/mman.h>
3  #include <sys/stat.h>
4  #include <sys/wait.h>
5  #include <unistd.h>
6  #include <fcntl.h>
7  #include <semaphore.h>
8  #include <signal.h>
9  #include <stdio.h>
10 #include <stdlib.h>
11 #include <errno.h>
12
13 #define SHARED_ARR "shared_array"
14 #define SEMAPHORE_FORKING "fork_semaphore"
15 #define N 6
16
17 /* to ensure that fp can be used by all processes */
18 static int fp;
19
20 void print_message(int signo) {
21     printf("%d got hit by signal %d\n", getpid(), signo);
22 }
23
24 void examine_child(int pid) {
25     int status;
26
27     if (waitpid(pid, &status, 0) == -1) {
28         perror("wait");
29         exit(errno);
30     }
31
32     if (WIFEXITED(status)) {
33         printf("%d exited with %d\n", pid, WEXITSTATUS(status));
34     }
35     else if (WIFSIGNALED(status)) {
36         printf("%d killed by signal %d\n", pid, WTERMSIG(status));
37     }
38     else {
39         printf("%d died in some other way\n", pid);
40     }
41 }
42
43 void *get_shared_arr() {
44     void *buf;
45     /* assert: fp is a file descriptor which points to the shared memory,
46      * and is inherited by all processes */
47     if ((buf = mmap(NULL, sizeof(int) * N, PROT_READ | PROT_WRITE, MAP_SHARED, fp, 0)) == MAP_FAILED) {
48         perror("shared memory mmap() failed");
49         exit(errno);
50     }
51     return buf;
52 }
53
54
55 int main(void) {
56
57     // PROCESS 0
58     // can see the pids of {1, 2}
59     sem_t *fork_sem, *pgid_sem, *exit_sem;
60     int *child_pid;
```

```

61     int ret;
62
63     if (signal(SIGUSR1, print_message) == SIG_ERR) {
64         perror("SIGUSR1");
65         return errno;
66     }
67
68     // link shared memory address to process- this will be visible in children
69     if ((child_pid= mmap(NULL, sizeof(int) * N,
70                          PROT_READ | PROT_WRITE,
71                          MAP_SHARED | MAP_ANONYMOUS,
72                          -1, 0)) == MAP_FAILED) {
73
74         perror("mmap");
75         return errno;
76     }
77
78     // create a shared memory map for semaphore
79     if ((fork_sem = mmap(NULL, sizeof(sem_t),
80                          PROT_READ | PROT_WRITE,
81                          MAP_SHARED | MAP_ANONYMOUS,
82                          -1, 0)) == MAP_FAILED) {
83
84         perror("semaphore");
85         return errno;
86     }
87     /* initialize semaphore with initial value 0- when 5 and 3 get
88      * created, this is incremented by 1 each. 0 will synchronize by calling down() on this twice */
89     if (sem_init(&fork_sem, 1, 0) == -1) {
90         perror("semaphore initialization");
91         return errno;
92     }
93
94     // create a shared memory map for semaphore
95     if ((exit_sem = mmap(NULL, sizeof(sem_t),
96                          PROT_READ | PROT_WRITE,
97                          MAP_SHARED | MAP_ANONYMOUS,
98                          -1, 0)) == MAP_FAILED) {
99
100         perror("semaphore");
101         return errno;
102     }
103
104     /* initialize semaphore with initial value 0- when signals have been sent,
105      * this will be incremented 5 times, wherein all processes will exit */
106     if (sem_init(&exit_sem, 1, 0) == -1) {
107         perror("semaphore initialization");
108         return errno;
109     }
110
111     // create a shared memory map for semaphore
112     if ((pgid_sem = mmap(NULL, sizeof(sem_t),
113                          PROT_READ | PROT_WRITE,
114                          MAP_SHARED | MAP_ANONYMOUS,
115                          -1, 0)) == MAP_FAILED) {
116
117         perror("semaphore");
118         return errno;
119     }
120
121     /* initialize semaphore with initial value 0- when all processes have
122      * called setpgid, the parent can send signals */
123     if (sem_init(&pgid_sem, 1, 0) == -1) {
124         perror("semaphore initialization");
125         return errno;
126     }
127

```

```

128     /* assert: now each child will have access to the semaphore, unless the
129     * memory region is purposely unliked */
130
131
132     if ((ret = fork()) == -1) {
133         perror("fork 1");
134         return errno;
135     }
136     else if (!ret) {
137         // CHILD 1
138         // can see pids of {5}
139         // child_pid = (int *)get_shared_arr();
140
141         if ((ret = fork()) == -1) {
142             perror("fork 5");
143             return errno;
144         }
145         else if (!ret) {
146             // CHILD 5
147             // can see pids of {}
148             // child_pid = (int *)get_shared_arr();
149
150             // wait for the PID of 3 to be available
151             if (sem_wait(fork_sem) == -1) {
152                 perror("P operation in 5");
153             }
154
155             // set own process group to 3
156             if (setpgid(0, child_pid[3]) == -1) {
157                 perror("setpgid(5, 3)");
158                 return errno;
159             }
160
161             if (sem_post(pgid_sem) == -1) {
162                 perror("setpgid(5, 3) done synchronization");
163                 return errno;
164             }
165
166             // wait for 0 to allow exiting
167             if (sem_wait(exit_sem) == -1) {
168                 perror("V operation in 4");
169                 return errno;
170             }
171
172             printf("%d is child of %d\n", getpid(), getppid());
173
174             return 0;
175         }
176         child_pid[5] = ret;
177
178
179
180
181         // wait for 0 to allow exiting
182         if (sem_wait(exit_sem) == -1) {
183             perror("V operation in 2");
184             return errno;
185         }
186
187         printf("%d is child of %d\n", getpid(), getppid());
188
189         examine_child(child_pid[5]);
190
191         return 0;
192     }
193     child_pid[1] = ret;
194

```

```

195     if ((ret = fork()) == -1) {
196         perror("fork 2");
197         return errno;
198     }
199     else if (!ret) {
200         // CHILD 5
201         // can see pids of {3}
202         // child_pid = (int *)get_shared_arr();
203
204         if ((ret = fork()) == -1) {
205             perror("fork 3");
206             return errno;
207         }
208         else if (!ret) {
209             // CHILD 3
210             // can see pids of {4}
211
212             child_pid[3] = getpid();
213
214             if (setpgid(child_pid[3], child_pid[3]) == -1) {
215                 perror("setpgid(3, 3)");
216                 return errno;
217             }
218
219             if (sem_post(pgid_sem) == -1) {
220                 perror("setpgid(3, 3) done synchronization");
221                 return errno;
222             }
223
224             // NOWWWW tell 5 that pid of 3 is available in shared memory, and it can call setpgid safely
225             if (sem_post(fork_sem) == -1) {
226                 perror("P operation in 3");
227                 return errno;
228             }
229
230
231             if ((ret = fork()) == -1) {
232                 perror("fork 4");
233                 return errno;
234             }
235             else if (!ret) {
236                 // CHILD 4
237                 // can see pids of {}
238                 // child_pid = (int *)get_shared_arr();
239
240                 if (setpgid(0, child_pid[3]) == -1) {
241                     perror("setpgid(4, 3)");
242                     return errno;
243                 }
244
245                 // NOWWW tell 0 that 4 has moved to new process group
246
247                 if (sem_post(pgid_sem) == -1) {
248                     perror("setpgid(4, 3) done synchronization");
249                     return errno;
250                 }
251
252
253                 // wait for 0 to allow exiting
254                 if (sem_wait(exit_sem) == -1) {
255                     perror("V operation in 4");
256                     return errno;
257                 }
258
259                 printf("%d is child of %d\n", getpid(), getppid());
260
261                 return 0;

```

```

262     }
263
264     child_pid[4] = ret;
265
266
267     if (sem_wait(exit_sem) == -1) {
268         perror("V operation in 3");
269         return errno;
270     }
271
272     printf("%d is child of %d\n", getpid(), getppid());
273
274     examine_child(child_pid[4]);
275
276     return 0;
277 }
278 child_pid[3] = ret;
279
280
281     // wait for 0 to allow exiting
282     if (sem_wait(exit_sem) == -1) {
283         perror("V operation in 2");
284         return errno;
285     }
286
287     printf("%d is child of %d\n", getpid(), getppid());
288
289     examine_child(child_pid[3]);
290
291     return 0;
292 }
293 child_pid[2] = ret;
294
295     // wait till all processes created, and the 3 setpgid calls finish
296     int i;
297     for (i = 0; i < 3; i++) {
298         if (sem_wait(pgid_sem) == -1) {
299             perror("V operation in 0");
300             return errno;
301         }
302     }
303
304     child_pid[0] = getpid();
305
306     if (kill(-child_pid[3], SIGUSR1) == -1) {
307         perror("kill");
308         return errno;
309     }
310
311     for (i = 0; i <= 5; i++) {
312         printf("Process %d = %d\n", i, child_pid[i]);
313     }
314
315
316     for (i = 0; i < 5; i++) {
317         if (sem_post(exit_sem) == -1) {
318             perror("P by 0");
319             return errno;
320         }
321     }
322
323     examine_child(child_pid[1]);
324     examine_child(child_pid[2]);
325
326     return 0;
327
328

```


Explanation

- Shared memory is created for sharing pids and semaphores
- NO. of Semaphores used = 3
- *fork_sem* is used by process 3 to tell process 5 that its pid is available in shared memory, and that its process group has been created
- *pgid_sem* is used for telling 0 that all 3 setpgid calls are done
- *exit_sem* is used for telling 1-5 that it has called kill, and they are free to exit now

Output

```

hp@aditi:~/Desktop/BTech/AUP/LAB/7$ gcc 3.c -o 3 -lpthread
hp@aditi:~/Desktop/BTech/AUP/LAB/7$ ./3
Process 0 = 58382
Process 1 = 58383
Process 2 = 58384
Process 3 = 58386
Process 4 = 58387
Process 5 = 58385
58383 is child of 58382
58384 is child of 58382
58387 got hit by signal 10
58385 got hit by signal 10
58385 is child of 58383
58387 is child of 58386
58385 exited with 0
58383 exited with 0
58386 got hit by signal 10
58386 is child of 58384
58387 exited with 0
58386 exited with 0
58384 exited with 0
hp@aditi:~/Desktop/BTech/AUP/LAB/7$

```

Figure 2: Execution of 3, SIGUSR hits 3, 4, 5