PCAP : Demonstration of MPI & openmp programs

Aditi Rajendra Medhane 111803177

1 : Hello World Non-blocking

Code

```c
#include <stdio.h>
#include <mpi.h>
#include <math.h>
#include <string.h>

int main(int argc, char* argv) {
    MPI_Status status;
    int num;

    MPI_Request request;
    //Initialize MPI computation
    MPI_Init(NULL, NULL);
    //Determine a process's ID number
    MPI_Comm_rank(MPI_COMM_WORLD, &num);

    double d = 100.0;
    char arr[] = "Hello World";

    int tag = 1;

    if(num == 0) {
        MPI_Isend(arr, strlen(arr)+1, MPI_BYTE, 1, tag, MPI_COMM_WORLD, &request);
        MPI_Irecv(arr, strlen(arr)+1, MPI_BYTE, 1, tag, MPI_COMM_WORLD, &request);
        MPI_Wait(&request, &status);
        printf("%s revcieved from process %d\n", arr, num);
    }
    else {
        MPI_Isend(arr, strlen(arr)+1, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &request);
        MPI_Irecv(arr, strlen(arr)+1, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &request);
        MPI_Wait(&request, &status);
        printf("%s revcieved from process %d\n", arr, num);
    }

    MPI_Finalize();
    return 0;
}
```

Output

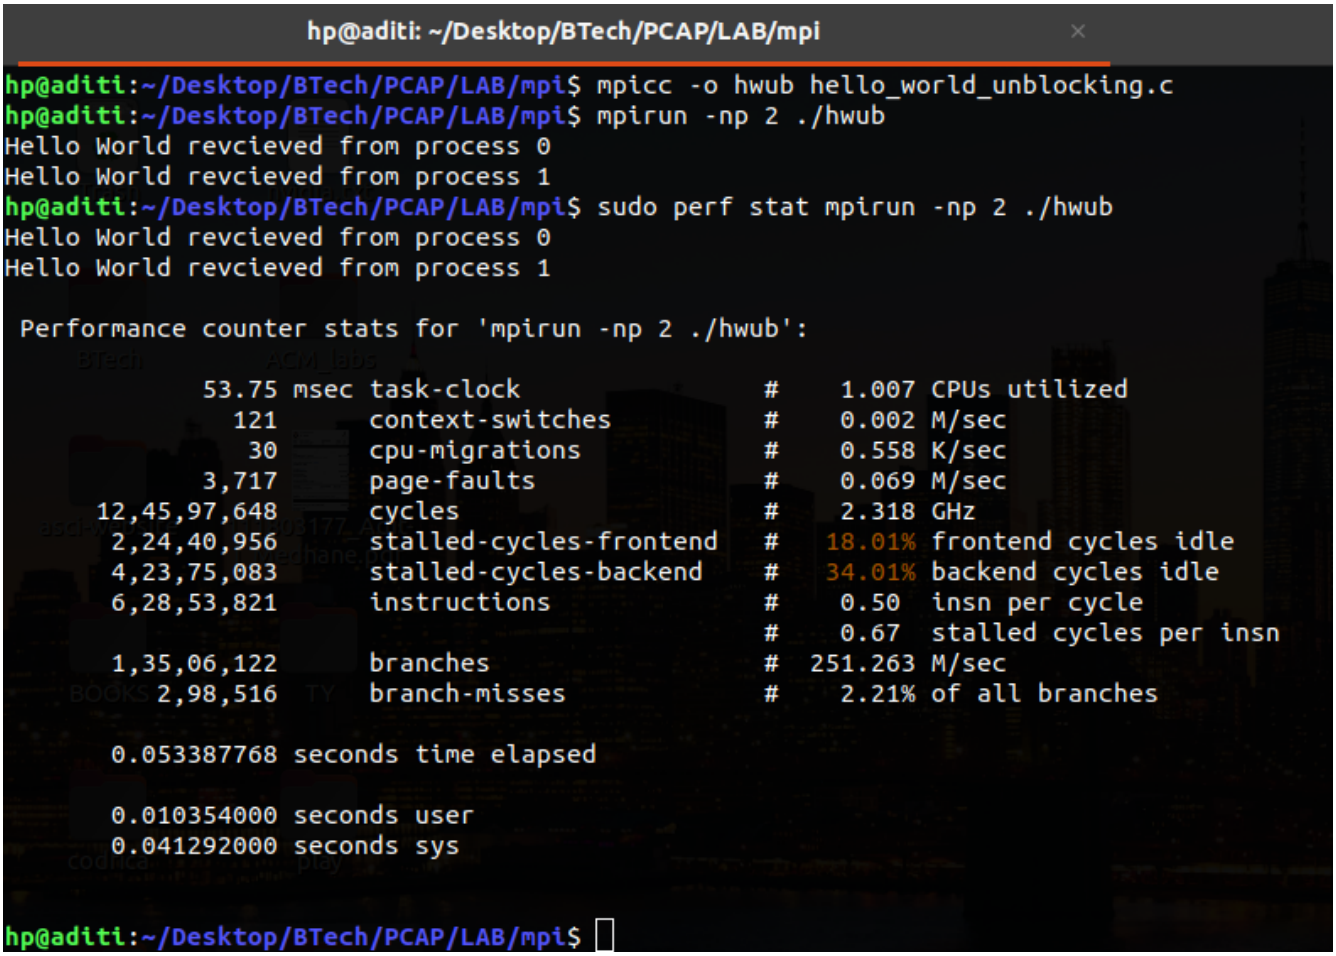

Figure 1: compilation, execution, perf stats of hello world non-blocking

2 : Hello World blocking

Code

```c
#include <stdio.h>
#include <mpi.h>
#include <math.h>
#include <string.h>

int main(int argc, char* argv) {
    MPI_Status status;
    int num;

    //Initialize MPI computation
    MPI_Init(NULL, NULL);
    //Determine a process's ID number
    MPI_Comm_rank(MPI_COMM_WORLD, &num);

    double d = 100.0;
    char arr[] = "Hello World";

    int tag = 1;

    if(num == 0) {
        MPI_Send(arr, strlen(arr)+1, MPI_BYTE, 1, tag, MPI_COMM_WORLD);
        MPI_Recv(arr, strlen(arr)+1, MPI_BYTE, 1, tag, MPI_COMM_WORLD, &status);
        printf("%s revcieved from process %d\n", arr, num);
    }
    else {
        MPI_Send(arr, strlen(arr)+1, MPI_BYTE, 0, tag, MPI_COMM_WORLD);
        MPI_Recv(arr, strlen(arr)+1, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &status);
        printf("%s revcieved from process %d\n", arr, num);
    }

    MPI_Finalize();
    return 0;
}
```

Output

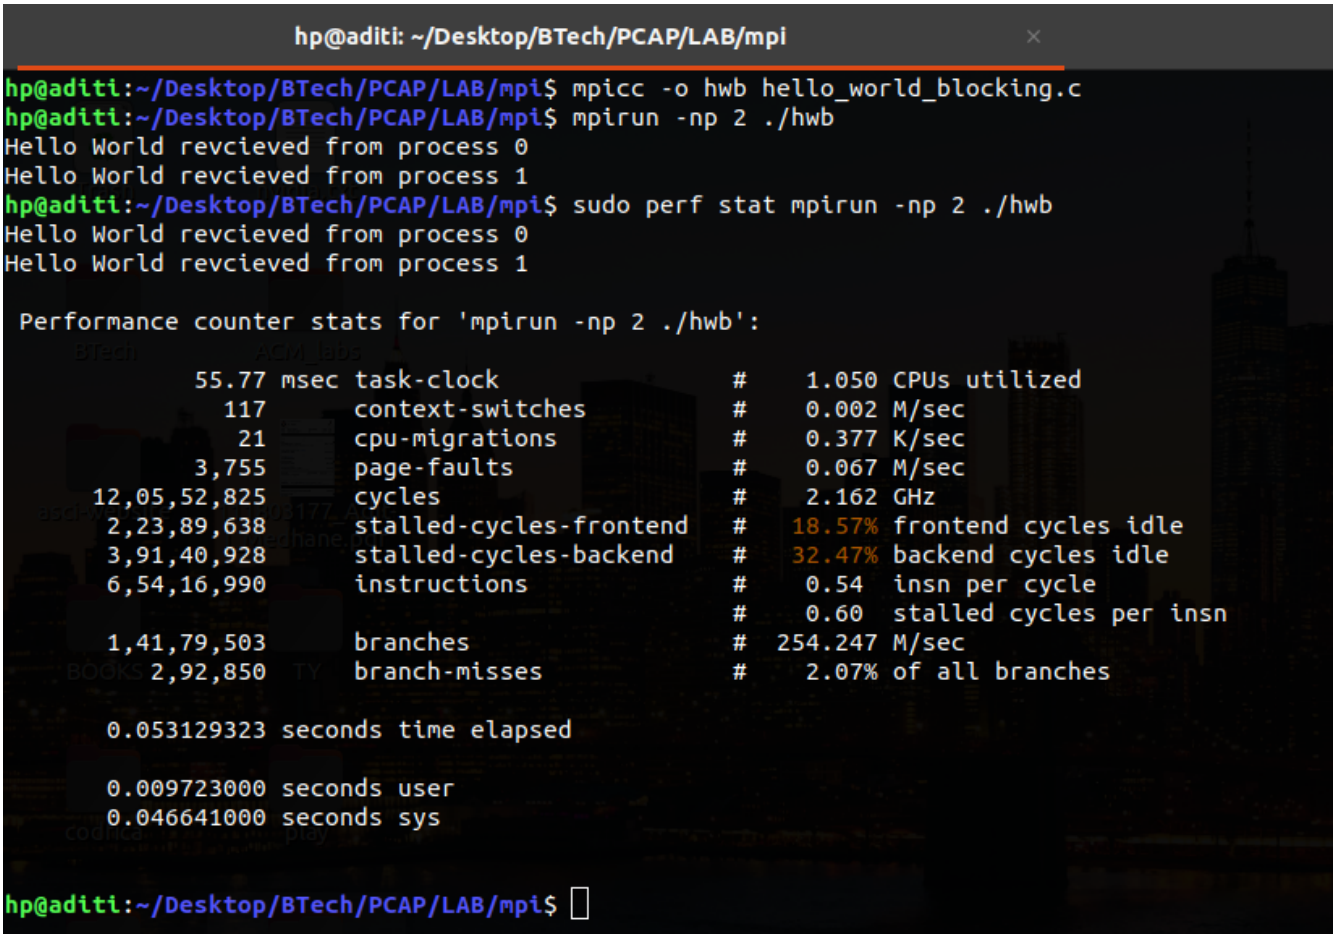

Figure 2: compilation, execution, perf stats of hello world blocking

3 : Trapezium rule Non-blocking

Implements the trapezoidal rule for numerical integration. To approximate the area between the graph of a function, y = f(x), two vertical lines, and the x-axis.

Explanation

Steps cosidered while parallelizing code : 1. Partition the problem solution into tasks. 2. Identify the communication channels between the tasks. 3. Aggregate the tasks into composite tasks. 4. Map the composite tasks to cores.

For the trapezoidal rule, two types of tasks will need to consider : 1. finding the area of a single trapezoid. 2. computing the sum of these areas. Then the communication channels will join each of the tasks of the first type to the single task of the second type.

Code

```c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <mpi.h>

const double a = 0;
const double b = 10000;

double trapezoid_area(double left_endpt, double right_endpt, int trap_count, double base_len);
double F(double x);


int main(int argc, char** argv) {
    int rank, size, n_trapezoids, n;
    double x0, x1, h, process_integral, final_integral;
    int source;

    MPI_Request request;
    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (argc!= 2){
        printf("Enter the command as : mpirun -np <N> %s <number of trapezoids> \n", argv[0]);
        n_trapezoids = -1;
        MPI_Finalize();
        exit(-1);
    }
    else {
        n_trapezoids = atoi(argv[1]);
    }
    MPI_Bcast(&n_trapezoids, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);

    //For every process, h and n will be same
    h = (b-a)/n_trapezoids;
    n = n_trapezoids/size;

    //For calculating the interval of integration for each process
    x0 = a + rank * n * h;
    x1 = x0 + n * h;

    MPI_Barrier(MPI_COMM_WORLD);

    //calculate integral of each process
    process_integral = trapezoid_area(x0, x1, n, h);

    if (rank != 0) {
        MPI_Isend(&process_integral, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, &request);
    }
    else {
        final_integral = process_integral;
            for (source = 1; source < size; source++) {
```

```
            MPI_Status status;
            MPI_Irecv(&process_integral, 1, MPI_DOUBLE, source, 0, MPI_COMM_WORLD, &request);
            MPI_Wait(&request, &status);
            final_integral += process_integral;
            }

            printf("For n = %d trapezoids:\n", n_trapezoids);
            printf("Integration of x^2 from %0.2f to %0.2f = %f\n", a, b, final_integral);
    }

    MPI_Finalize();

    return 0;

}

double F(double x) {
    return x * x;
}

double trapezoid_area(double left_endpt, double right_endpt, int trapezoid_count, double base_len) {
    double integral, x;
    int i;

    integral = (F(left_endpt) + F(right_endpt))/2.0;
    for (i = 1; i <= trapezoid_count-1; i++) {
        x = left_endpt + i * base_len;
        integral += F(x);
    }
    integral = integral * base_len;

    return integral;
}
```

Output

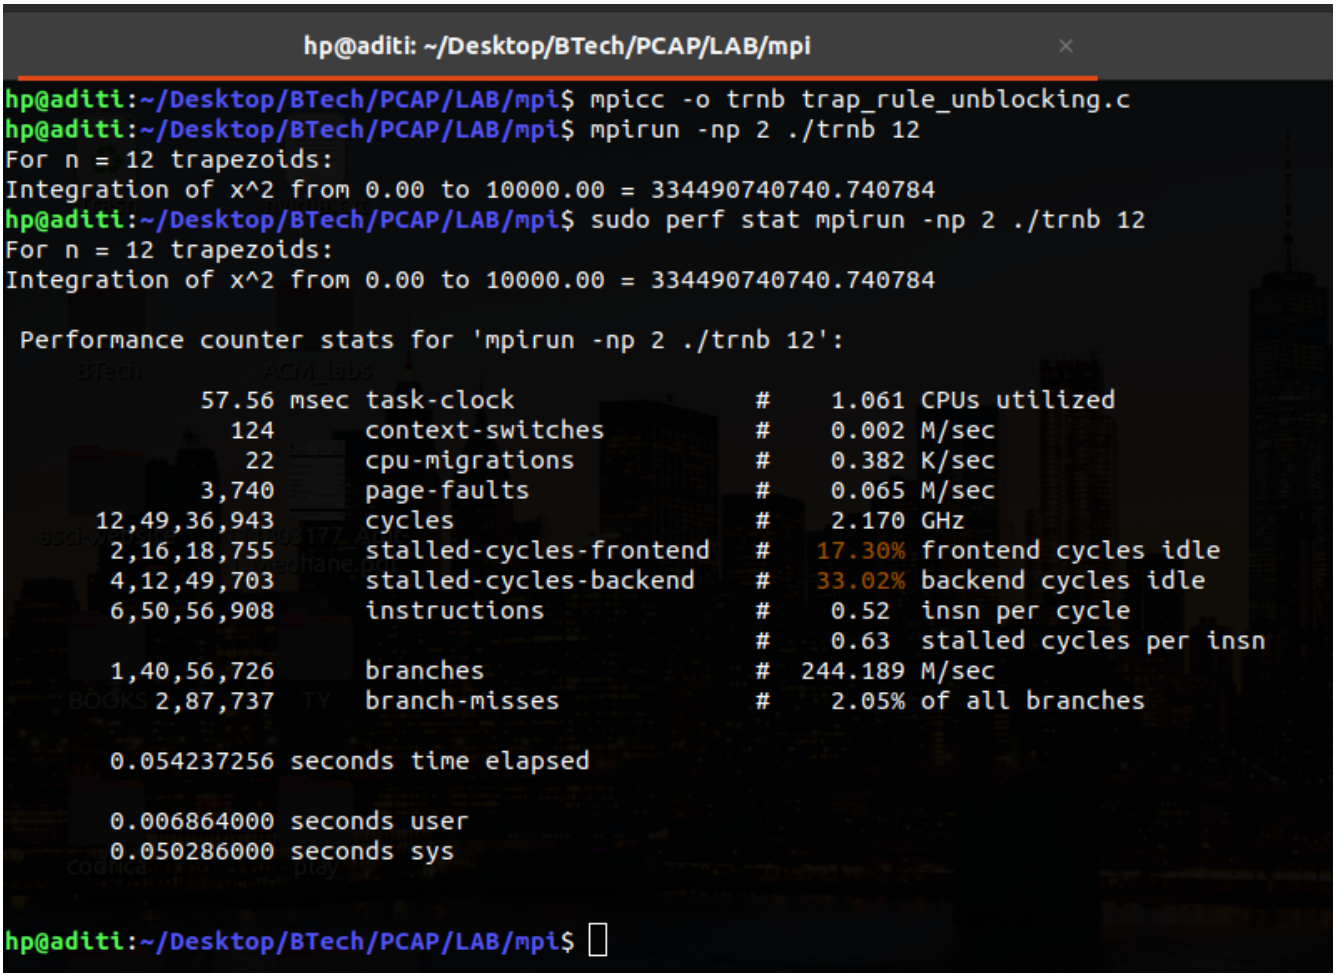

Figure 3: compilation, execution, perf stats of trapezium non-blocking

4 : Trapezium rule blocking

Code

```c
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <mpi.h>

const double a = 0;
const double b = 10000;

double trapezoid_area(double left_endpt, double right_endpt, int trap_count, double base_len);
double F(double x);


int main(int argc, char** argv) {
    int rank, size, n_trapezoids, n;
    double x0, x1, h, process_integral, final_integral;
    int source;

    MPI_Init(NULL, NULL);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (argc!= 2){
        printf("Enter the command as : mpirun -np <N> %s <number of trapezoids> \n", argv[0]);
        n_trapezoids = -1;
        MPI_Finalize();
        exit(-1);
    }
    else {
        n_trapezoids = atoi(argv[1]);
    }
    MPI_Bcast(&n_trapezoids, 1, MPI_DOUBLE, 0, MPI_COMM_WORLD);

    //For every process, h and n will be same
    h = (b-a)/n_trapezoids;
    n = n_trapezoids/size;

    //For calculating the interval of integration for each process
    x0 = a + rank * n * h;
    x1 = x0 + n * h;

    MPI_Barrier(MPI_COMM_WORLD);

    //calculate integral of each process
    process_integral = trapezoid_area(x0, x1, n, h);

    if (rank != 0) {
        MPI_Send(&process_integral, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD);
    }
    else {
        final_integral = process_integral;
            for (source = 1; source < size; source++) {
            MPI_Status status;
            MPI_Recv(&process_integral, 1, MPI_DOUBLE, source, 0, MPI_COMM_WORLD, &status);
            final_integral += process_integral;
            }

            printf("For n = %d trapezoids:\n", n_trapezoids);
            printf("Integration of x^2 from %0.2f to %0.2f = %f\n", a, b, final_integral);
    }

    MPI_Finalize();

    return 0;
```

```c

}

double F(double x) {
    return x * x;
}

double trapezoid_area(double left_endpt, double right_endpt, int trapezoid_count, double base_len) {
    double integral, x;
    int i;

    integral = (F(left_endpt) + F(right_endpt))/2.0;
    for (i = 1; i <= trapezoid_count-1; i++) {
        x = left_endpt + i * base_len;
        integral += F(x);
    }
    integral = integral * base_len;

    return integral;
}
```

Output

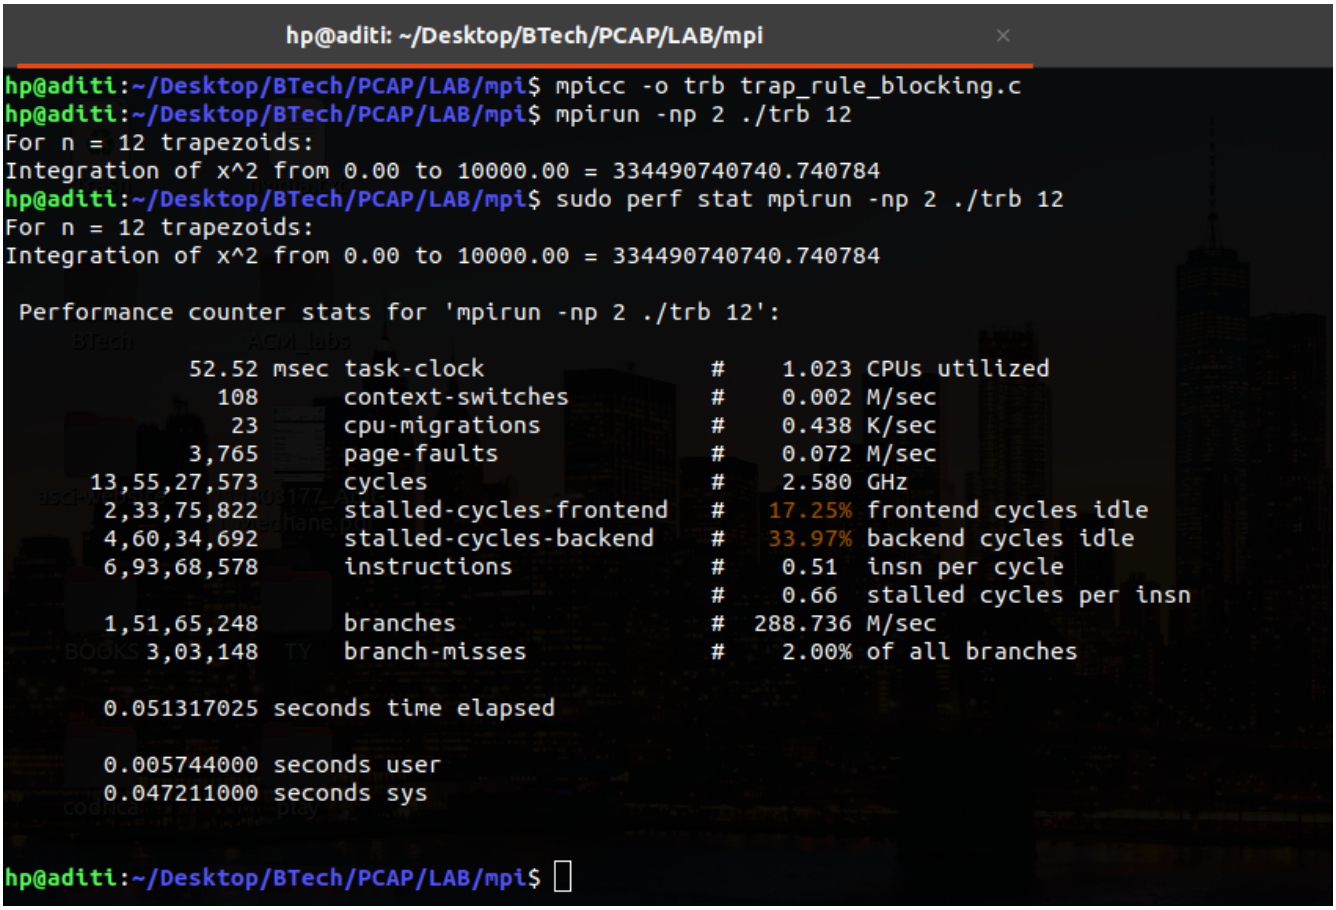

Figure 4: compilation, execution, perf stats of trapezium blocking

5 : Matrix Multiplication using MPI

Code

```c
#include <stdio.h>
#include <mpi.h>
#include <math.h>
#include<unistd.h>
#include<stdlib.h>

MPI_Status status;


int main(int argc, char **argv)  {
    int rank, size;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if(argc < 2) {
        printf("Enter the correct arguments\n");
        return 1;
    }

    int N = atoi(argv[1]);

    double a[N][N],b[N][N],c[N][N];


    int workers, rows, offset, dest, source, r1=N, c1=N,i,j,k;

    if(rank == 0)  {
        for (i = 0; i < r1; i++) {
            for (j = 0; j < c1; j++) {
                a[i][j] = rand() % 10;
                b[i][j] = rand() % 10;
            }
        }

        workers = size - 1;


        rows = N/workers;
        offset = 0;

        for(dest = 1; dest <= workers; dest++) {
            MPI_Send(&offset, 1, MPI_INT, dest, 1, MPI_COMM_WORLD);
            MPI_Send(&rows, 1, MPI_INT, dest, 1, MPI_COMM_WORLD);
            MPI_Send(&a[offset][0], rows*N, MPI_DOUBLE,dest,1, MPI_COMM_WORLD);
            MPI_Send(&b, N*N, MPI_DOUBLE, dest, 1, MPI_COMM_WORLD);
            offset = offset + rows;
        }

        for (i=1; i<=workers; i++) {
            source = i;
            MPI_Recv(&offset, 1, MPI_INT, source, 2, MPI_COMM_WORLD, &status);
            MPI_Recv(&rows, 1, MPI_INT, source, 2, MPI_COMM_WORLD, &status);
            MPI_Recv(&c[offset][0], rows*N, MPI_DOUBLE, source, 2, MPI_COMM_WORLD, &status);
        }


        printf("Matrix multiplication is done\n");
        /*
        for (i=0; i<N; i++) {
            for (j=0; j<N; j++)
                printf("%6.2f   ", c[i][j]);
```

```
            printf ("\n");
        }
        */
    }

    if (rank > 0) {
        source = 0;
        MPI_Recv(&offset, 1, MPI_INT, source, 1, MPI_COMM_WORLD, &status);
        MPI_Recv(&rows, 1, MPI_INT, source, 1, MPI_COMM_WORLD, &status);
        MPI_Recv(&a, rows*N, MPI_DOUBLE, source, 1, MPI_COMM_WORLD, &status);
        MPI_Recv(&b, N*N, MPI_DOUBLE, source, 1, MPI_COMM_WORLD, &status);

        for (k=0; k<N; k++) {
            for (i=0; i<rows; i++) {
                c[i][k] = 0.0;
                for (j=0; j<N; j++)
                    c[i][k] = c[i][k] + a[i][j] * b[j][k];
                }
        }

        MPI_Send(&offset, 1, MPI_INT, 0, 2, MPI_COMM_WORLD);
        MPI_Send(&rows, 1, MPI_INT, 0, 2, MPI_COMM_WORLD);
        MPI_Send(&c, rows*N, MPI_DOUBLE, 0, 2, MPI_COMM_WORLD);

    }

    MPI_Finalize();
    return 0;
}
```

Conclusion

Size | Time Elapsed(MPI) | CPU Utilization(MPI) |

128 0.069979909 seconds 1.208 CPUs utilized

256 0.324119474 seconds 1.802 CPUs utilized

512 6.599078551 seconds 1.945 CPUs utilized

Output

```
hp@aditi:~/Desktop/BTech/PCAP/LAB/mpi$ time sudo perf stat mpirun -np 2 ./mul 128
Matrix multiplication is done

 Performance counter stats for 'mpirun -np 2 ./mul 128':

             83.79 msec task-clock                #    1.208 CPUs utilized
                122      context-switches          #    0.001 M/sec
                 19      cpu-migrations            #    0.227 K/sec
              3,970      page-faults               #    0.047 M/sec
     23,17,44,889      cycles                    #    2.766 GHz
      4,11,47,286      stalled-cycles-frontend   #   17.76% frontend cycles idle
      6,94,86,010      stalled-cycles-backend    #   29.98% backend cycles idle
     25,65,65,752      instructions              #    1.11  insn per cycle
                                                 #    0.27  stalled cycles per insn
      4,08,69,135      branches                  #  487.781 M/sec
         3,29,605      branch-misses             #    0.81% of all branches

       0.069334002 seconds time elapsed

       0.052938000 seconds user
       0.030064000 seconds sys


real    0m0.102s
user    0m0.071s
sys     0m0.039s
```

```
hp@aditi:~/Desktop/BTech/PCAP/LAB/mpi$ time sudo perf stat mpirun -np 2 ./mul 256
Matrix multiplication is done

 Performance counter stats for 'mpirun -np 2 ./mul 256':

            584.08 msec task-clock                #    1.802 CPUs utilized
               260      context-switches          #    0.445 K/sec
                10      cpu-migrations            #    0.017 K/sec
             4,585      page-faults               #    0.008 M/sec
    2,01,96,02,558      cycles                    #    3.458 GHz
      29,16,96,399      stalled-cycles-frontend   #   14.44% frontend cycles idle
      77,59,60,967      stalled-cycles-backend    #   38.42% backend cycles idle
    2,52,92,64,919      instructions              #    1.25  insn per cycle
                                                  #    0.31  stalled cycles per insn
      47,24,37,359      branches                  #  808.856 M/sec
          4,55,553      branch-misses             #    0.10% of all branches

       0.324119474 seconds time elapsed

       0.519555000 seconds user
       0.063155000 seconds sys


real    0m0.373s
user    0m0.527s
sys     0m0.085s
```

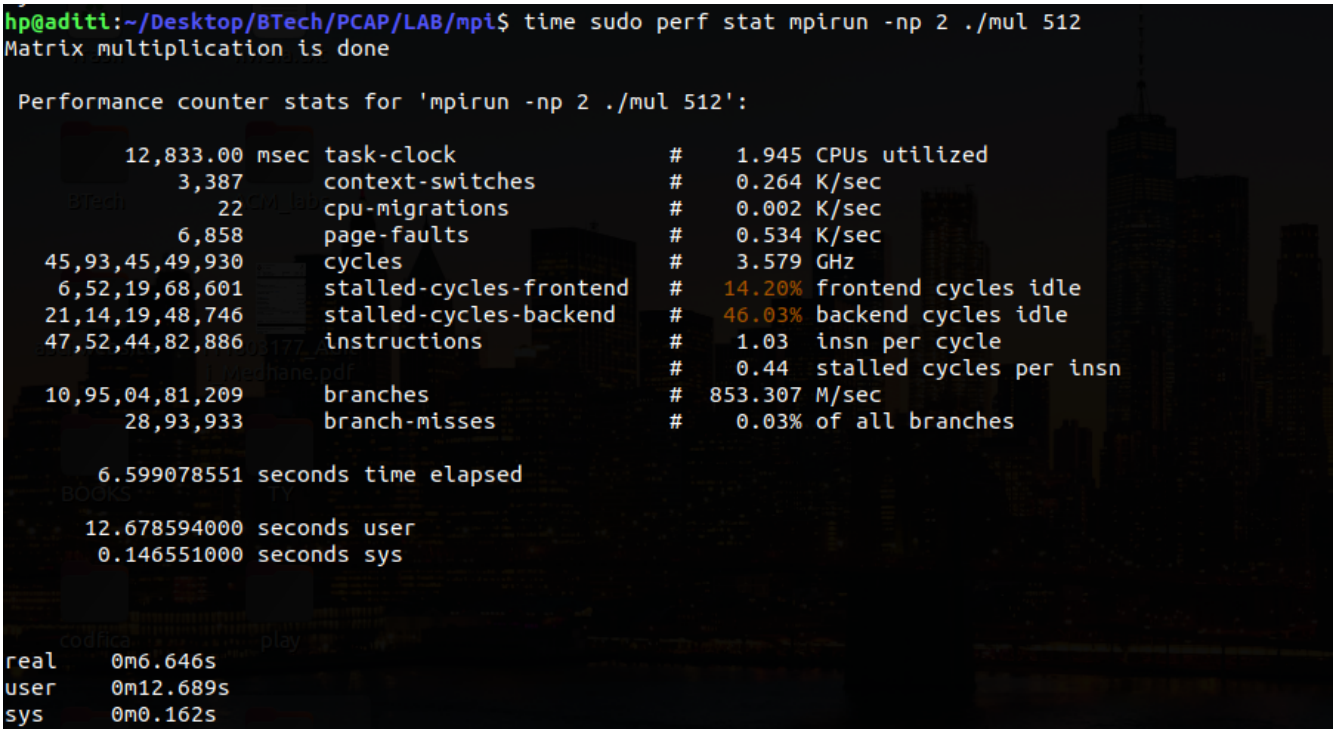
hp@aditi:~/Desktop/BTech/PCAP/LAB/mpi$ time sudo perf stat mpirun -np 2 ./mul 512
Matrix multiplication is done
Performance counter stats for 'mpirun -np 2 ./mul 512':
12,833.00 msec task-clock # 1.945 CPUs utilized
3,387 context-switches # 0.264 K/sec
22 cpu-migrations # 0.002 K/sec
6,858 page-faults # 0.534 K/sec
45,93,45,49,930 cycles # 3.579 GHz
6,52,19,68,601 stalled-cycles-frontend # 14.20% frontend cycles idle
21,14,19,48,746 stalled-cycles-backend # 46.03% backend cycles idle
47,52,44,82,886 instructions # 1.03 insn per cycle
0.44 stalled cycles per insn
10,95,04,81,209 branches # 853.307 M/sec
28,93,933 branch-misses # 0.03% of all branches
6.599078551 seconds time elapsed
12.678594000 seconds user
0.146551000 seconds sys
real 0m6.646s
user 0m12.689s
sys 0m0.162s

6 : Matrix Multiplication using openmp

Code

```c
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <omp.h>

#define THREADS 16

int main(int argc, char **argv) {
    // Initialize the matrices
    int SIZE = atoi(argv[1]);

    int matrix_a[SIZE][SIZE];
    int matrix_b[SIZE][SIZE];
    int matrix_c[SIZE][SIZE];

    for(int i = 0; i < SIZE; i++) {
        for(int j = 0; j < SIZE; j++) {
            matrix_a[i][j] = rand() % 10;
            matrix_b[i][j] = rand() % 10;
            matrix_c[i][j] = 0;
        }
    }

    omp_set_num_threads(THREADS);
    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        #pragma omp for
            for(int i = id*(SIZE/THREADS); i < (id + 1)*(SIZE/THREADS); i++) {
                for(int c = 0; c < SIZE; c++) {
                    matrix_c[i][c] = 0;
                    for(int k = 0; k < SIZE; k++) {
                        matrix_c[i][c] += matrix_a[i][k]*matrix_b[k][c];
                    }
                }
            }
    }

    return 0;
}
```

Observation

Size | Time Elapsed 128 0.003665319 seconds

256 0.009580259 seconds

512 0.137304077 seconds

Output

Conclusion

MPI => Internodes Openmp => Intranodes

MPI : Runtime decreases as the numeber of core increases, upto a limit where there is not much improvement.

Openmp : Using, more threads decreases the run time, the desired result, but for small matrix sizes using more number of threads slow down the calculation.

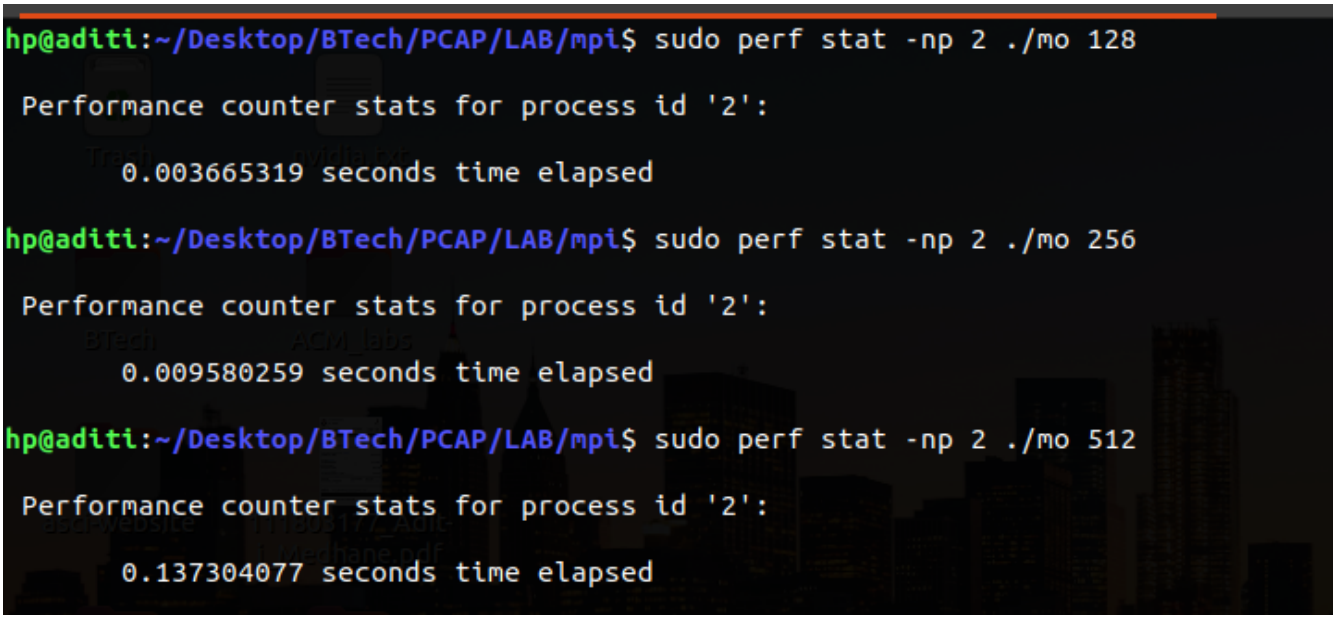

Figure 5: Output